

1. Preprocessing For Text Sentiment Analysis

Preprocessing is the process of taking the raw text data and converting it into a numerical format that a machine learning model, like a BERT model, can understand and learn from.

The preprocessing in your code involves three main stages:

1. Data Reduction & Mapping (4-Class Problem)
 2. Label Encoding
 3. Tokenization
-

1. Data Reduction & Mapping (GoEmotions ---> 4 Classes)

The original GoEmotions dataset is a multi-label classification problem with 27 distinct emotion classes. This is too complex for this specific goal (mental state classification), so the first step simplifies the problem.

- Goal: Convert the complex 27 original emotion classes into just 4 simplified, broader classes for a more focused task (like mental health-related classification).
 - The Mapping: A custom dictionary (`map_manual`) is used to group the 27 original classes into the new 4 classes:
 - Original: *sadness, disappointment, grief* ---> New: Depression
 - Original: *fear, nervousness, panic, surprise* ---> New: Anxiety
 - Original: *anger, annoyance, disgust* ---> New: Frustration
 - Original: *joy, relief, neutral, love, optimism* (and all others) ---> New: Calmness
 - Final Data Frame: The code loops through every text example, applies this mapping, and creates new, clean dataframes (train, validation, test) where each text only has one final label out of the new 4 classes.
-

2. Label Encoding (Text Labels ---> Numbers)

While we have the 4 new simplified *text labels* (depression, anxiety, frustration, calmness), the model ultimately needs numbers to work with.

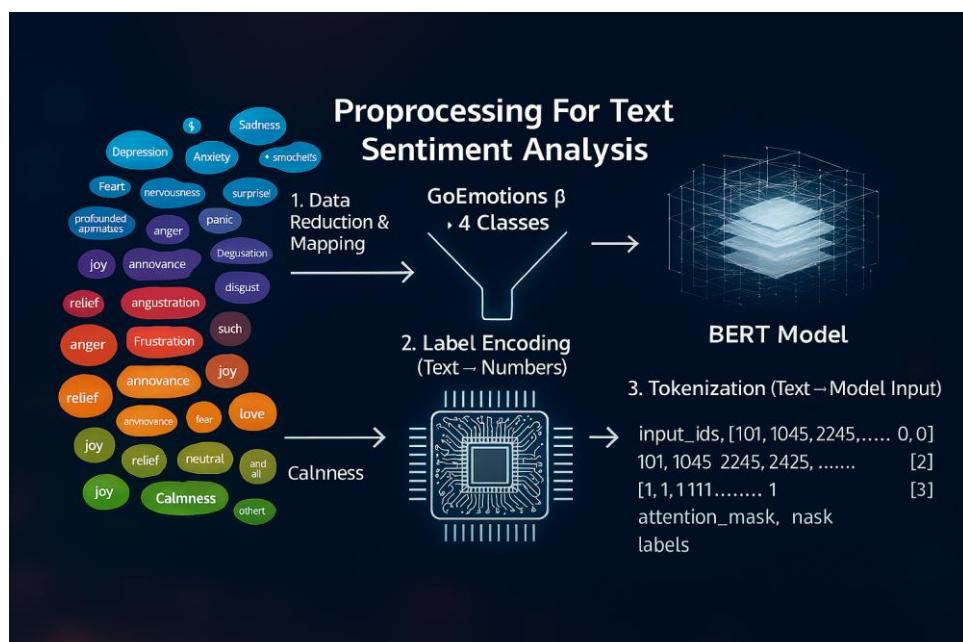
- Goal: Convert the text label for each example into a unique numerical ID.
- The Encoding: A dictionary (`LABEL2ID`) assigns an integer to each class:
 - Depression ---> 0
 - Anxiety ---> 1

- Frustration ---> 2
 - Calmness ---> 3
 - Result: A new column, `label_id`, is added to the data, which is what the model will learn to predict.
-

3. Tokenization (Text ---> Model Input)

This is the most crucial step that prepares the raw text into the specific numerical format required by the BERT model.

- Goal: Turn the raw text into a list of numbers (tokens) that the BERT tokenizer understands.
- The Process:
 1. Load Tokenizer: The code loads a specific tokenizer (e.g., "bert-base-uncased").
 2. Conversion: For every text example:
 - It breaks the text into "tokens" (words or sub-words).
 - It converts each token into a corresponding ID number (e.g., "The" ---> 1996). This creates the `input_ids` list.
 - It adds a special `attention_mask`, which is a list of 1s and 0s telling the model which tokens are *real words* (1) and which are just *padding* (0).
 3. Standardization: All resulting lists are standardized to the same length (`max_length=128`) by either truncating long texts or padding short ones with zeros.
- Final Output: The datasets are now ready for training, containing the necessary columns: `input_ids`, `attention_mask`, and the numerical labels (which was the old `label_id`).



2. Preprocessing For Facial Data Analysis

For an image-based model, preprocessing involves preparing the image files themselves and mapping their labels.

The preprocessing in your face analysis code involves three main stages:

1. Data Reduction & Mapping (7-Class ---> 4-Class)
 2. Image Scaling & Standardization
 3. Data Augmentation (For Robustness)
-

1. Data Reduction & Mapping (7 ---> 4 Classes)

The original FER2013 dataset has 7 standard emotion classes. Just like with the text data, we simplify this to the 4 broader classes needed for your specific analysis.

- Goal: Convert the 7 original emotion classes (like 'angry', 'fear') into the 4 simplified classes (like 'frustration', 'anxiety').
 - The Mapping: A custom dictionary (`emotion_map`) groups the original classes:
 - Original: angry, disgust ---> New: Frustration
 - Original: fear, surprise ---> New: Anxiety
 - Original: sad ---> New: Depression
 - Original: happy, neutral ---> New: Calmness
 - Implementation: The raw data generator sees the original 7-class structure, but a custom function (`custom_generator`) applies this mapping on-the-fly to every batch, converting the 7-class numerical label into the new 4-class numerical label.
-

2. Image Scaling & Standardization

Before the images can be fed into the model, their pixels (which are initially numbers between 0 and 255) must be standardized.

- Goal: Ensure all images are uniform in size and have their pixel values scaled down, which helps the model learn faster and better.
- Resizing & Grayscale: The `ImageDataGenerator` is instructed to:
 - Resize every image to a small, standard size (48x48 pixels).
 - Convert the images to grayscale (`color_mode="grayscale"`), meaning each pixel only has one value (intensity) instead of three (RGB).

3. Data Augmentation (For Robustness)

Data augmentation is a technique used only on the training data to artificially create more training examples and make the final model more robust.

- **Goal:** Teach the model to correctly identify an emotion even if the face is slightly tilted, shifted, or zoomed differently than in the original training photo.
- **The Process:** The train_datagen uses several techniques to slightly modify the images every time they are loaded:
 - **Rotation:** Tilts the face slightly (rotation_range=20).
 - **Shifts:** Moves the face horizontally or vertically (width_shift_range, height_shift_range).
 - **Zoom:** Zooms in or out slightly (zoom_range).
 - **Flip:** Mirrors the image (horizontal_flip=True).
- **Result:** The model trains on millions of variations of the original photos, preventing it from overfitting to specific image compositions and improving its ability to generalize to new, unseen faces.